



IZMIR

@dockerizmir

Build, Ship, Run


Docker is the world's leading software containerization platform



Demystifying Docker Swarm Mode

Ajeet Singh Raina

Docker Captain – Docker, Inc.



Who Am I?

- Sr. Systems Development Engineer at DellEMC
- 1st half of my career was in CGI & VMware
- 2nd half of my career has been in System Integration
- Testing/Project Lead for Dell EMC.
- Definitely more IT pro than developer
- @ajeetsraina (a frequent Twitterati)

The {code} Catalyst Program!



<http://www.collabnix.com>

Agenda

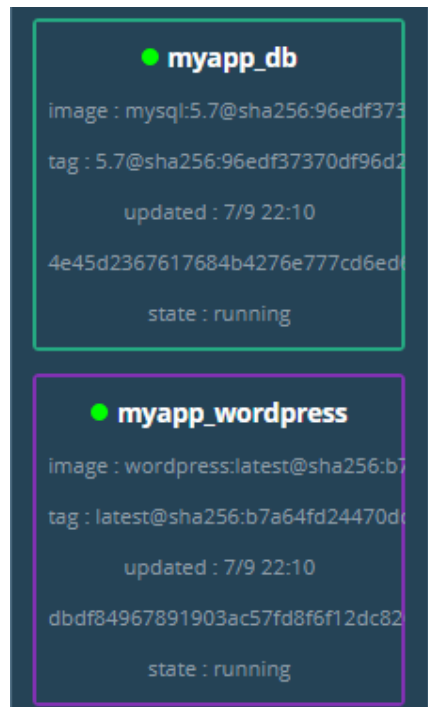
- Introduction to Docker Swarm
- Docker Swarm Mode Features
- Docker Stack Deployment
- What's new in Docker 17.06 Swarm Mode - Hybrid Swarm Setup, Topology Scheduling
- Demo – Hybrid Swarm(Play with Docker)

Introduction to Docker Swarm Mode



A Little Background: What is Swarm?

Let's start with Single Docker Host Application



A Docker Host

<http://collabnix.com/getting-started-with-docker-swarm/>


A Little Background: What is Swarm?

You want to add more hosts..

● manager1 manager 60.000G RAM	● manager2 manager 60.000G RAM	● manager3 manager 60.000G RAM	● worker1 worker 60.000G RAM	● worker2 worker 60.000G RAM
<div>● myapp_db image : mysql:5.7@sha256:96edf373 tag : 5.7@sha256:96edf37370df96d2 updated : 7/9 22:11 4b9cd79cad5d30f304300b10f35a4d state : running</div>	<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:9 2d43caa2953e423af3b9e4f756ded9 state : running</div>	<div>● myapp_db image : mysql:5.7@sha256:96edf373 tag : 5.7@sha256:96edf37370df96d2 updated : 7/9 22:12 9b75e2083ed7d6744f93d52ce9d742 state : running</div>	<div>● myapp_db image : mysql:5.7@sha256:96edf373 tag : 5.7@sha256:96edf37370df96d2 updated : 7/9 22:9 30d021d31fce5d19f3aa34d053d175 state : running</div>	<div>● myapp_db image : mysql:5.7@sha256:96edf373 tag : 5.7@sha256:96edf37370df96d2 updated : 7/9 22:10 a7d6e3bebcf9030992492843064f7 state : running</div>
<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:10 dbdf84967891903ac57fd8f6f12dc82 state : running</div>	<div>● myapp_db image : mysql:5.7@sha256:96edf373 tag : 5.7@sha256:96edf37370df96d2 updated : 7/9 22:10 732c68b848e2ef5ffa0fbc18633761 state : running</div>	<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:10 3a079a6b23ab0f26cc882d8ddfd3dd state : running</div>	<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:10 5091940001e66d60d290baa549fb state : running</div>	<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:10 a8bc43c6dc56910cc2b6df45f3225 state : running</div>
<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:10 d09df4513d70d1e6cbe97e7b3eb8 state : running</div>	<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:9 428877ab2a6d47c4fe60db186f219 state : running</div>	<div>● myapp_db image : mysql:5.7@sha256:96edf373 tag : 5.7@sha256:96edf37370df96d2 updated : 7/9 22:11 b64e90367aad63070f926a313e9513 state : running</div>	<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:10 a61396b0553b201962d5f6e1ac3613 state : running</div>	<div>● myapp_wordpress image : wordpress:latest@sha256:b7a64fd24470d tag : latest@sha256:b7a64fd24470d updated : 7/9 22:10 b04efc7983fe79747c0bb52d30aa1a state : running</div>

A Little Background: What is Swarm?

But Wait...



High Availability

Rolling Updates

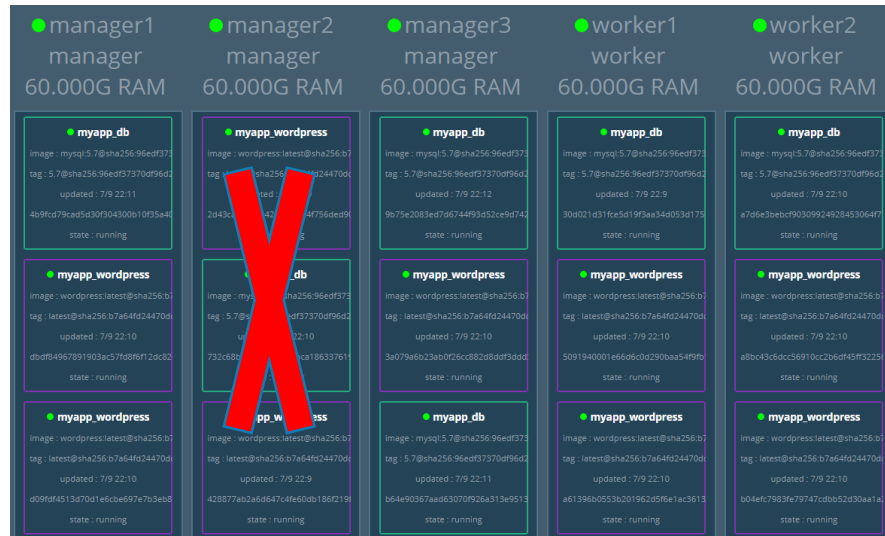
Service Discovery

Scalability

Scheduling

Failure Management

Container Security



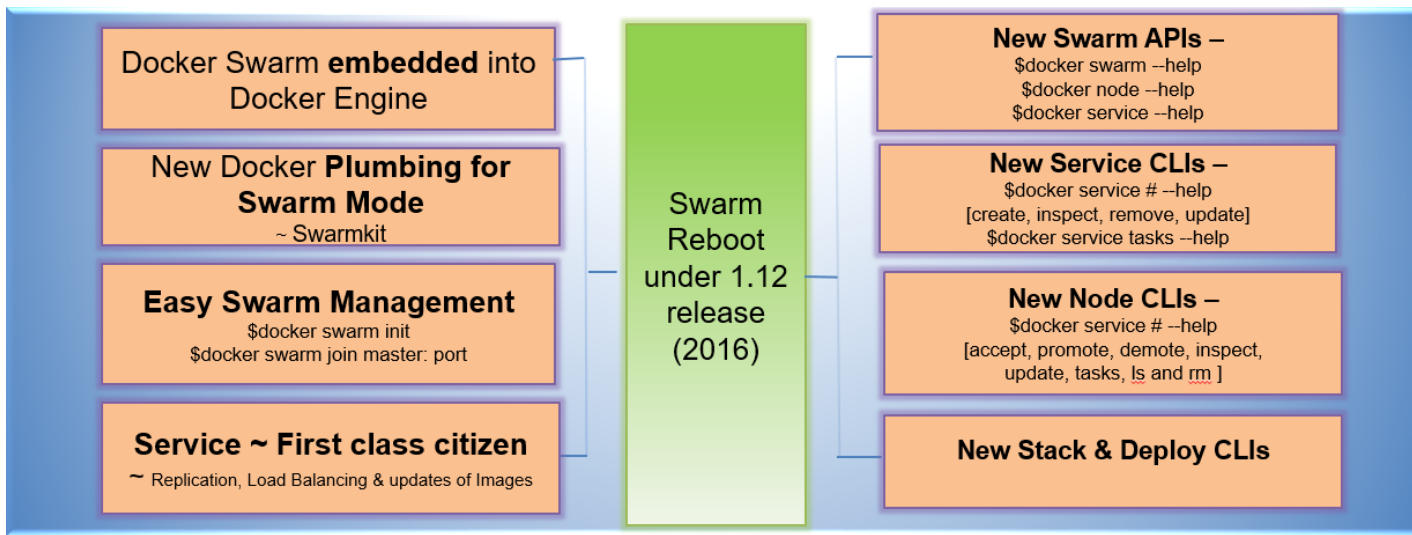
Docker Swarm Mode comes to rescue..



<http://collabnix.com/new-docker-1-12-comes-with-built-in-distribution-orchestration-system/>

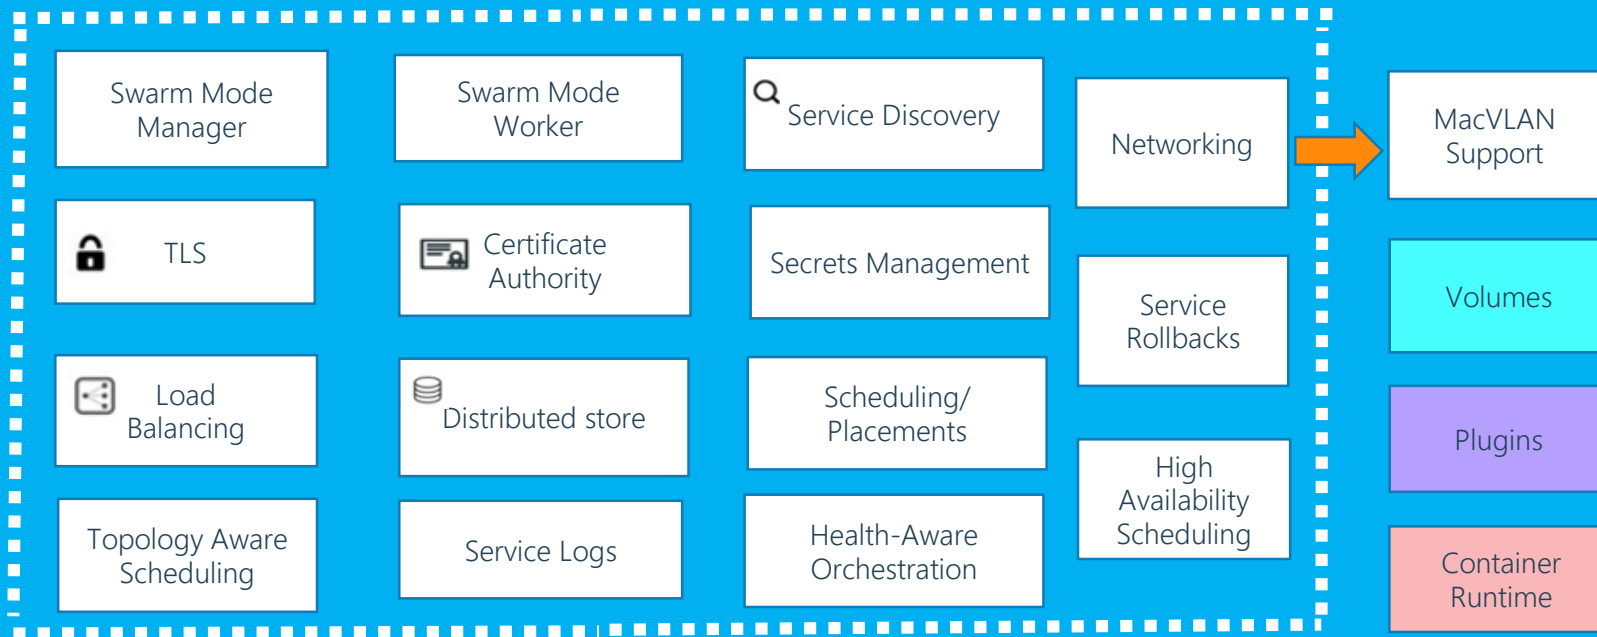
What is Swarm Mode?

- A swarm consists of one or more nodes: physical or virtual machines running Docker Engine.
- It was introduced first under Docker 1.12 release.
- It enables the ability to deploy containers across multiple Docker hosts, using overlay networks for service discovery with a built-in load balancer for scaling the services.



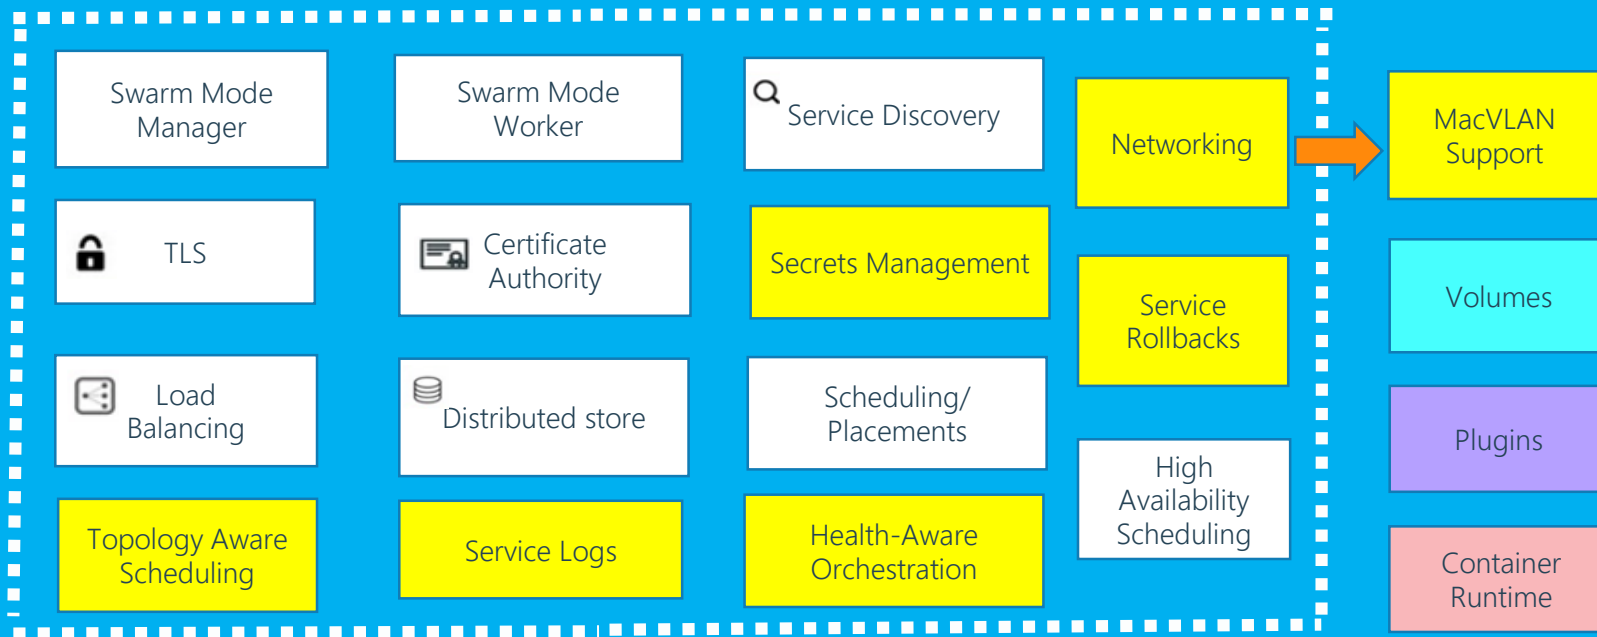
Swarm Mode Features under Docker 17.06

Orchestration Components



Swarm Mode Features under Docker 17.06

Orchestration Components



Building a Swarm Topology

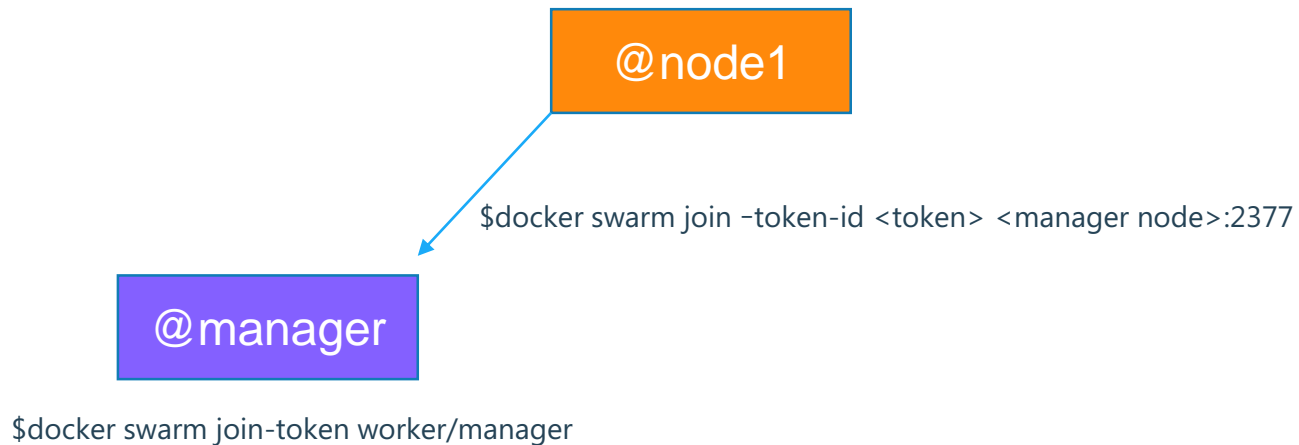


Building a Swarm Topology – Manual Way

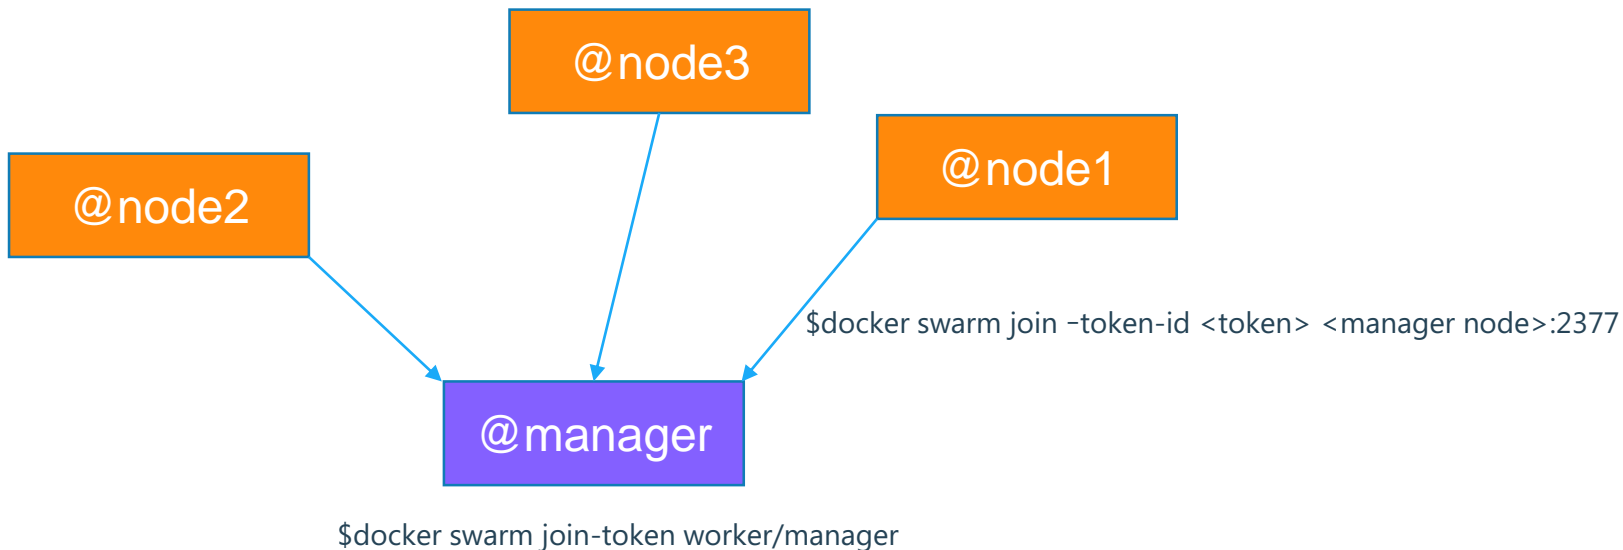
@manager

```
$docker swarm init --advertise-addr <IP of manager node>:2377  
or  
$docker swarm init --listen-addr <IP of manager node>:2377
```

Building a Swarm Topology – Manual Way



Building a Swarm Topology – Manual Way



```
$docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
1wmnidfmcop2uenypd1e8uh5k *	manager	Ready	Active	Leader
9ai9hbfm2ugkmjyt7kpx4hi9k	node1	Ready	Active	Reachable
j8iylu9npx5z1edepqp3shdw7	node2	Ready	Active	Reachable
fd8iylu9npx5z1edepqp3shdw7	node3	Ready	Active	Reachable

Building Swarm Topology – Scripted Method(Docker Machine)

```
master_node=master
node01_node=node01
node02_node=node02
```

```
# Initialize Virtualbox machines using Docker Machine
for i in $master_node $node01_node $node02_node; do docker-machine create -d virtualbox $i; done
```

```
# Create the Docker swarm
```

```
docker $(docker-machine config $master_node) swarm init \
--advertise-addr $(docker-machine ip $master_node):2377
```

@manager

```
# Add workers to the Docker Swarm
```

```
docker $(docker-machine config $node01_node) swarm join \
--token `docker $(docker-machine config $master_node) swarm join-token worker -q` \
$(docker-machine ip $master_node):2377
```

@node1,2

```
docker $(docker-machine config $node02_node) swarm join \
--token `docker $(docker-machine config $master_node) swarm join-token worker -q` \
$(docker-machine ip $master_node):2377
```

```
# Run the Docker Swarm visualizer
```

```
docker $(docker-machine config $master_node) run -it -d -p 5000:5000 \
-e HOST=`docker-machine ip $master_node` \
-e PORT=5000 \
-v /var/run/docker.sock:/var/run/docker.sock \
dockersamples/visualizer
```


Building a Swarm Topology – Cloud

Docker for AWS



Docker for Azure


Docker for GCP



Google Cloud Platform

Deployment manager

Building Swarm Topology – Hybrid Cluster

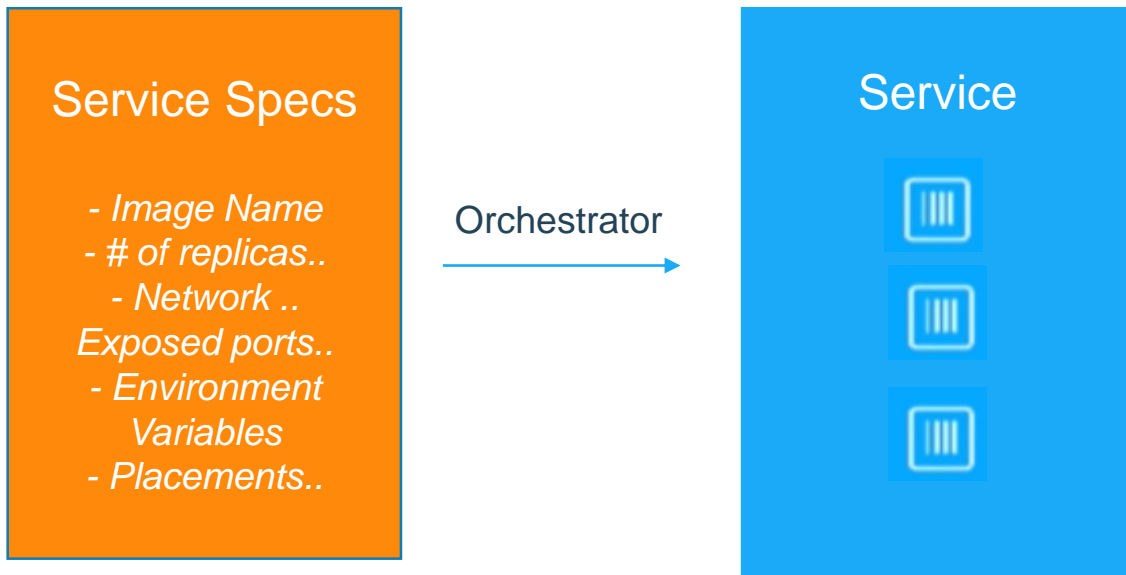


<http://collabnix.com/building-hybrid-docker-swarm-mode-cluster-on-google-cloud-platform/>

Service Discovery



Swarm is built on Services



What is Service?

- A definition of tasks to be executed on the worker nodes
- Central structure of swarm system
- An Evolution of `docker run` command
- It manages replicated set of containers
- A task carries a Docker container + commands to run inside the container.


• Container:

```
—docker run -d nginx
```

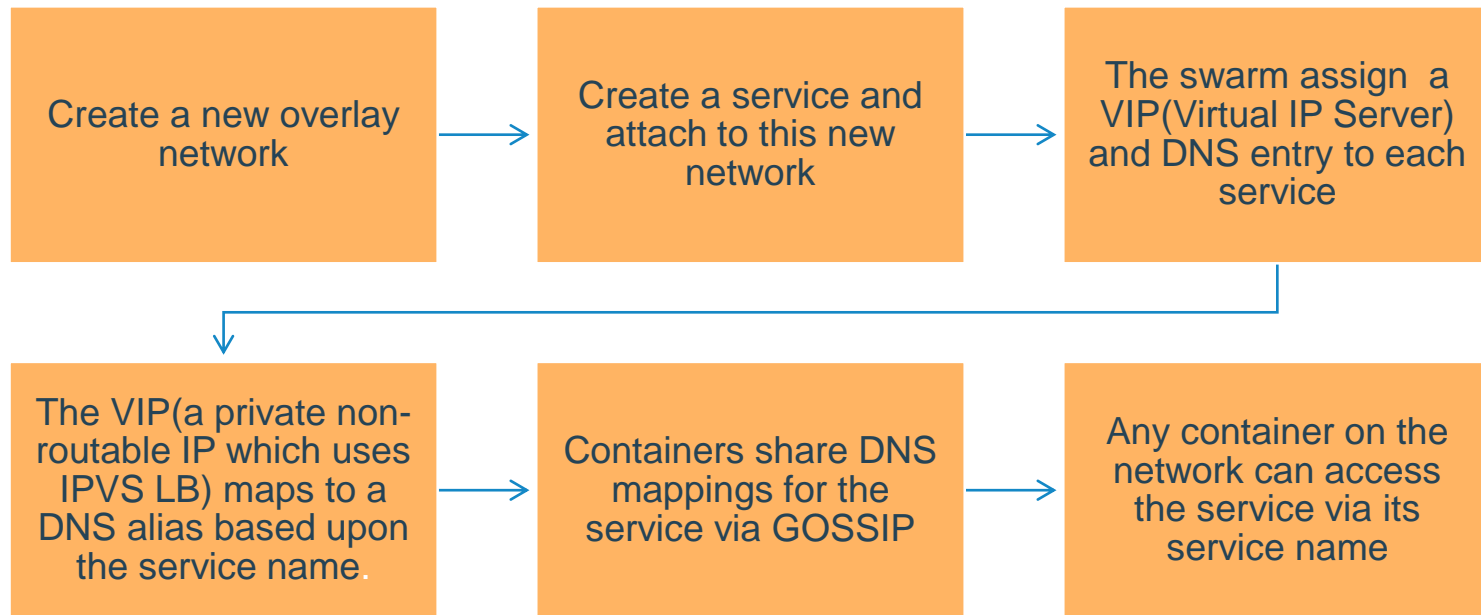


• Service:

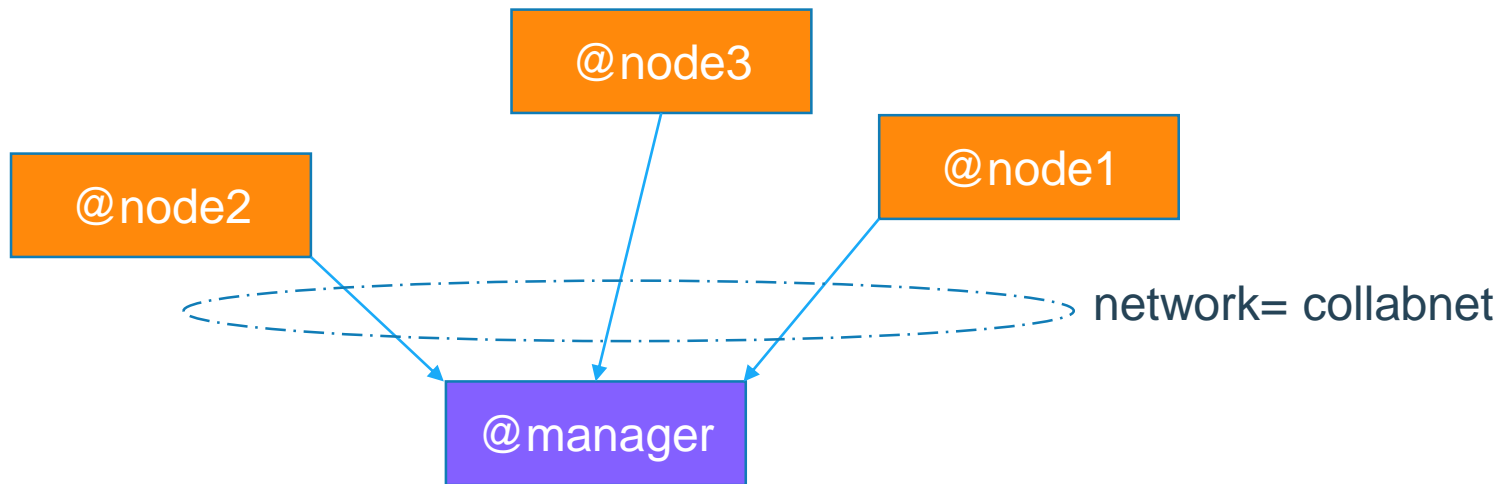
```
—docker service create --replicas 5 nginx
```



How Service Discovery works in Swarm Mode?



Building Our First Swarm Service



```
$docker network create -d overlay mynetwork
```

Swarm Cluster Setup

```
master==>sudo docker node ls
```

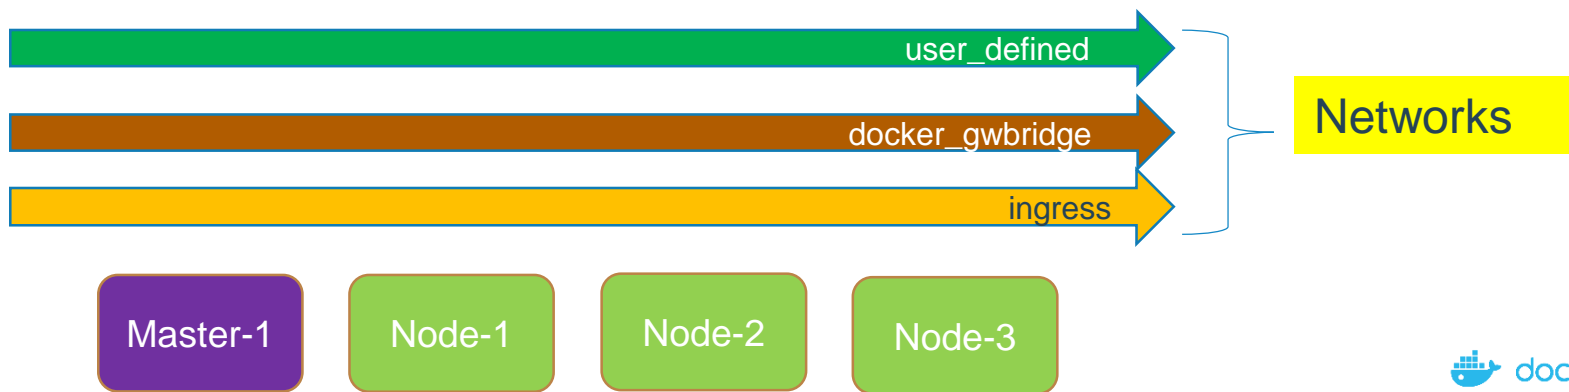
ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
1y2tfoe4jo6wjgyqizmmifkx5	* master1	Ready	Active	Leader
4rzd5krhbo5y2rr4axslrw4jm	test-node2	Ready	Active	
4w5oasgfglpdzngslrbalwqj6	test-node1	Ready	Active	
5t7ftnm6w44edvkgz9trw016	test-node3	Ready	Active	

```
root@master1:~# sudo docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
6f673c1acf3c	bridge	bridge	local
6eaeba0e7c6d	docker_gwbridge	bridge	local
84280982cb88	host	host	local
c4caizphmdpu	ingress	overlay	swarm
ac5edd465975	none	null	local

- The default gateway network
- The only network with connectivity to the outside world.(Port 7946 for network discovery)

- It is an overlay network on all exposed ports exist.
- Follows a node port model(each service has the same port on every node in the cluster).
- Numbered from 30000 through 32000.
- Used for Routing Mesh(Port 4789 for Ingress)

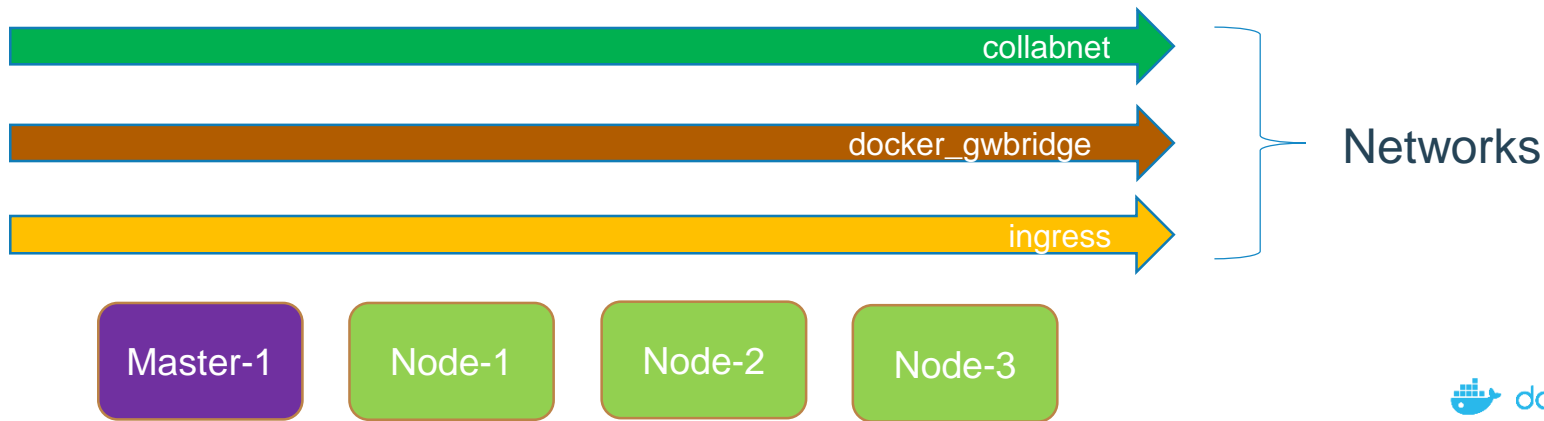


Creating a new overlay network

```
$ docker network create \
  --driver overlay \
  collabnet
```

```
master==>sudo docker network ls
```

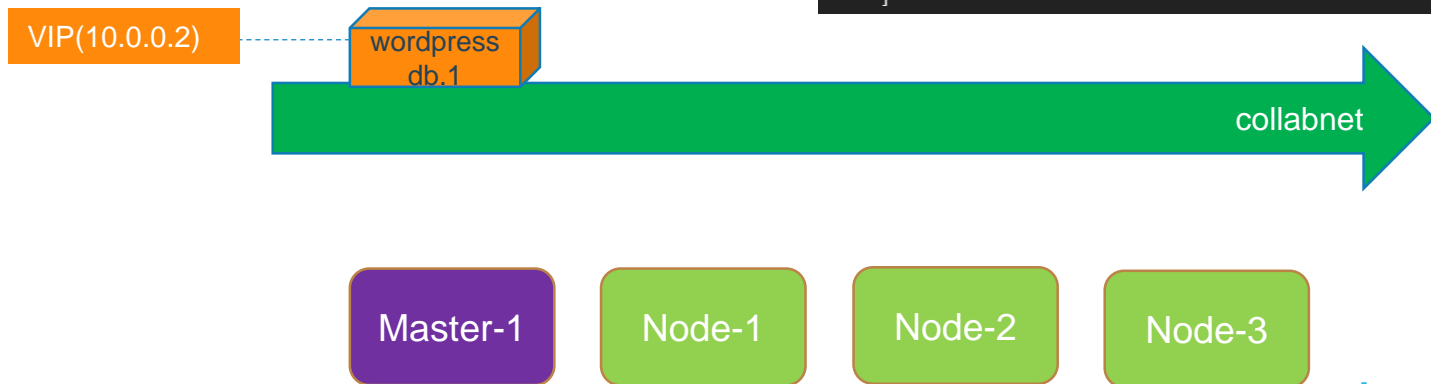
NETWORK ID	NAME	DRIVER	SCOPE
1f19d0986d1d	bridge	bridge	local
9eyjm4uv4ynm	collabnet	overlay	swarm
6eaeba0e7c6d	docker_gwbridge	bridge	local
84280982cb88	host	host	local
c4caizphmdpu	ingress	overlay	swarm
ac5edd465975	none	null	local



Creating a service "wordpressdb"

```
$ docker service create \
  --replicas 1 \
  --name wordpressdb \
  --network collabnet \
  --env MYSQL_ROOT_PASSWORD=collab123 \
  --env MYSQL_DATABASE=wordpress \
  --name wordpressdb \
  mysql:latest
```

```
{
  "Endpoint": {
    "Spec": {
      "Mode": "vip"
    },
    "VirtualIPs": [
      {
        "NetworkID": "9eyjm4uv4ynmz0aubfqxise29",
        "Addr": "10.0.0.2/24"
      }
    ]
  }
}
```



Creating a service "wordpressapp"

```
$ docker service create \
```

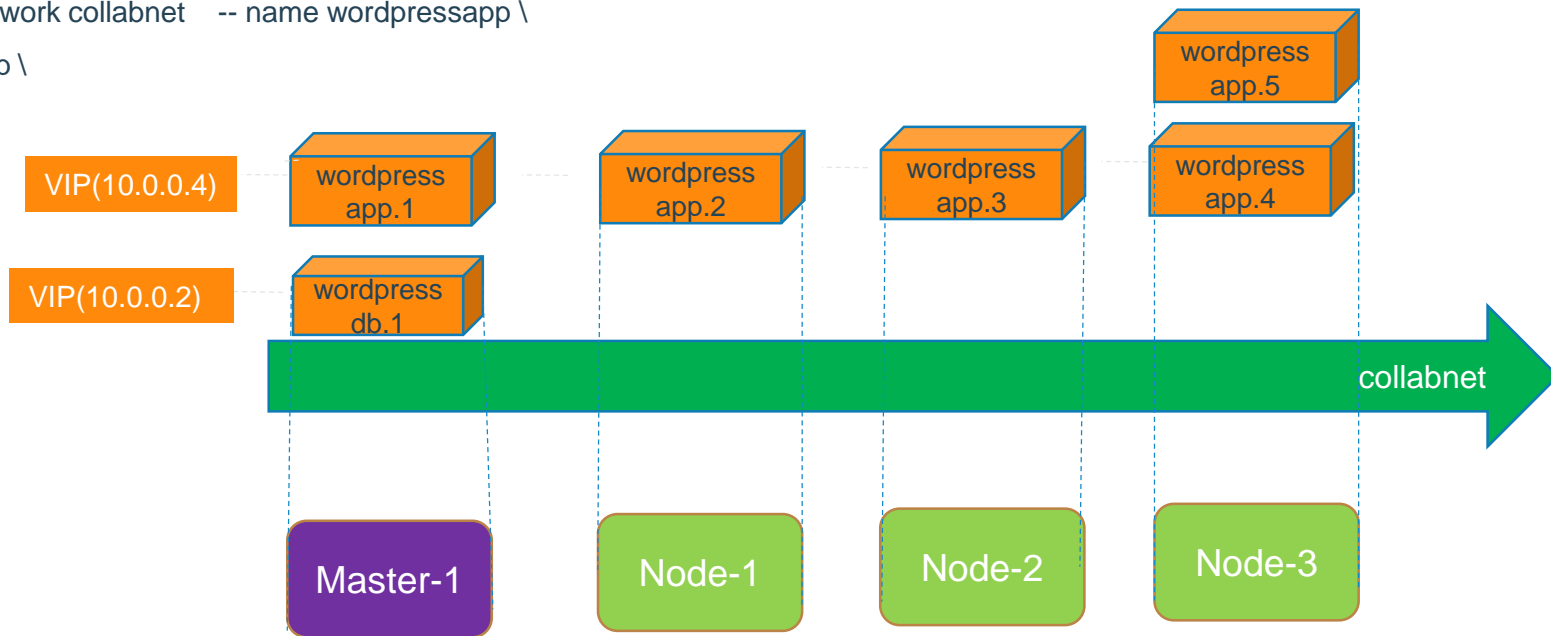
```
--env WORDPRESS_DB_HOST=wordpressdb \
```

```
--env WORDPRESS_DB_PASSWD=collab123 \
```

```
--replicas 5 --network collabnet -- name wordpressapp \
```

```
--publish 80:80/tcp \
```

```
wordpress:latest
```



Inspecting the services

```
$ docker service inspect \
--format=='{{json .Endpoint.VirtualIPs}}' \
wordpressapp
```

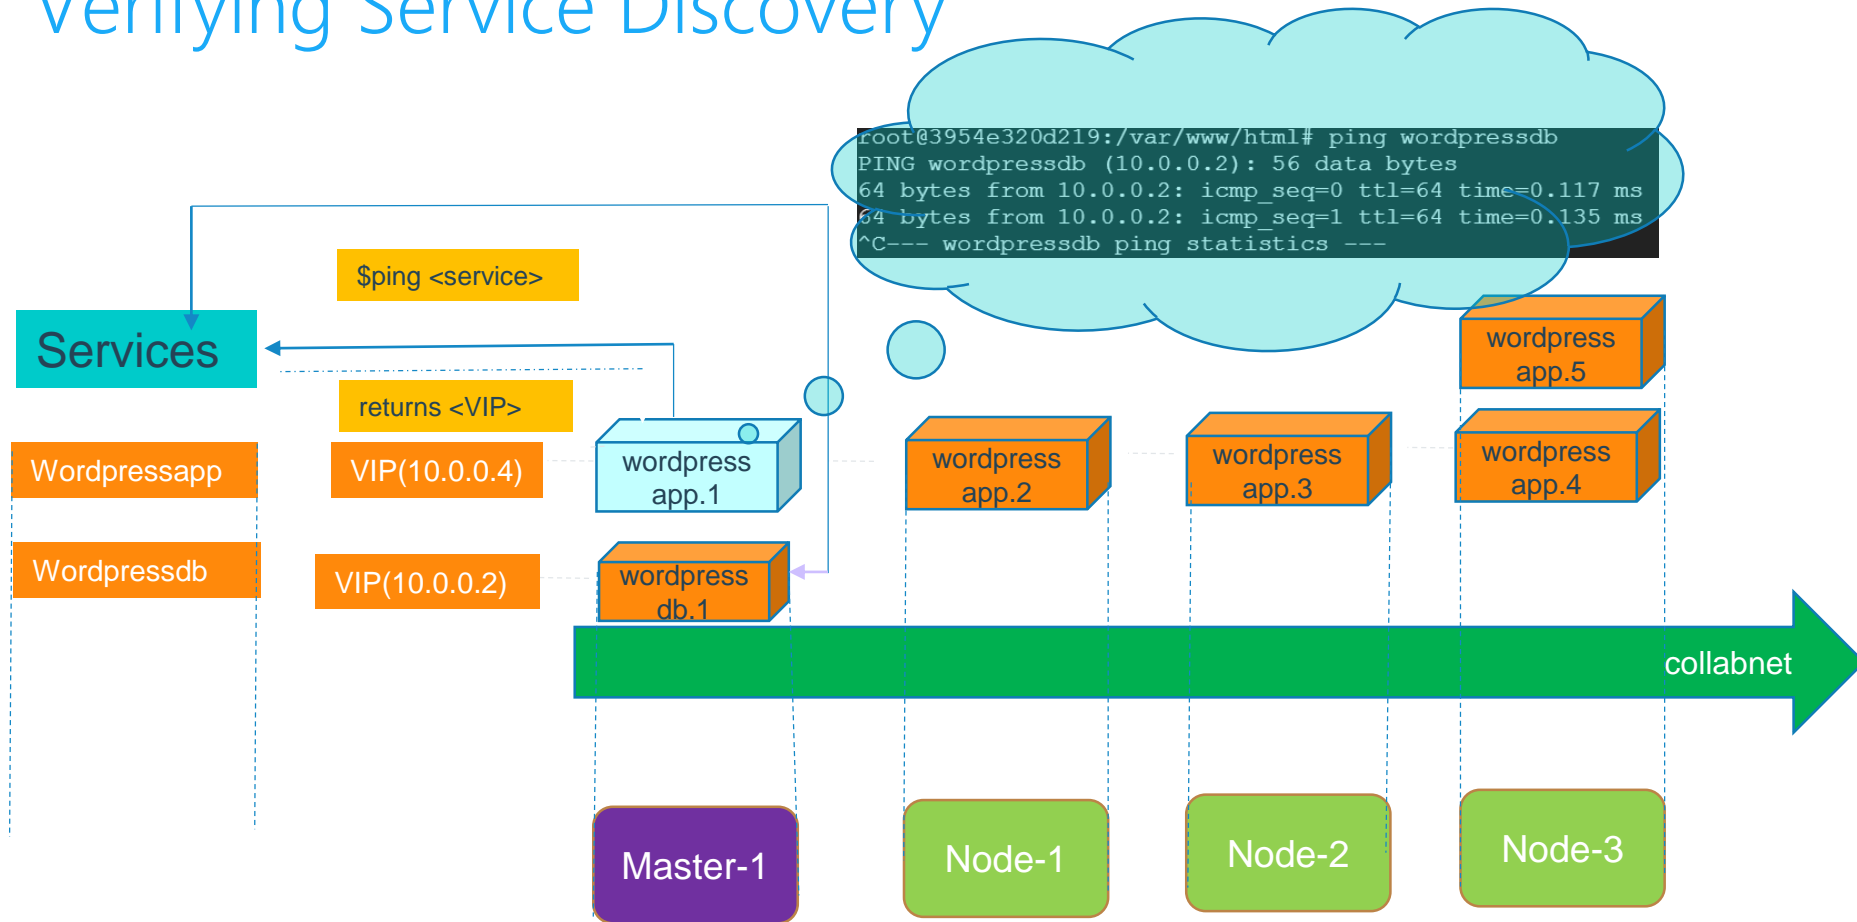
```
[{"NetworkID":"c4caizphmdpuhm1gjdle8eaal","Addr":"10.255.0.7/16"},
{"NetworkID":"9eyjm4uv4ynmz0aubfqxise29","Addr":"10.0.0.4/24"}]
```

```
$ docker service inspect \
--format=='{{json .Endpoint.VirtualIPs}}' \
wordpressdb
```

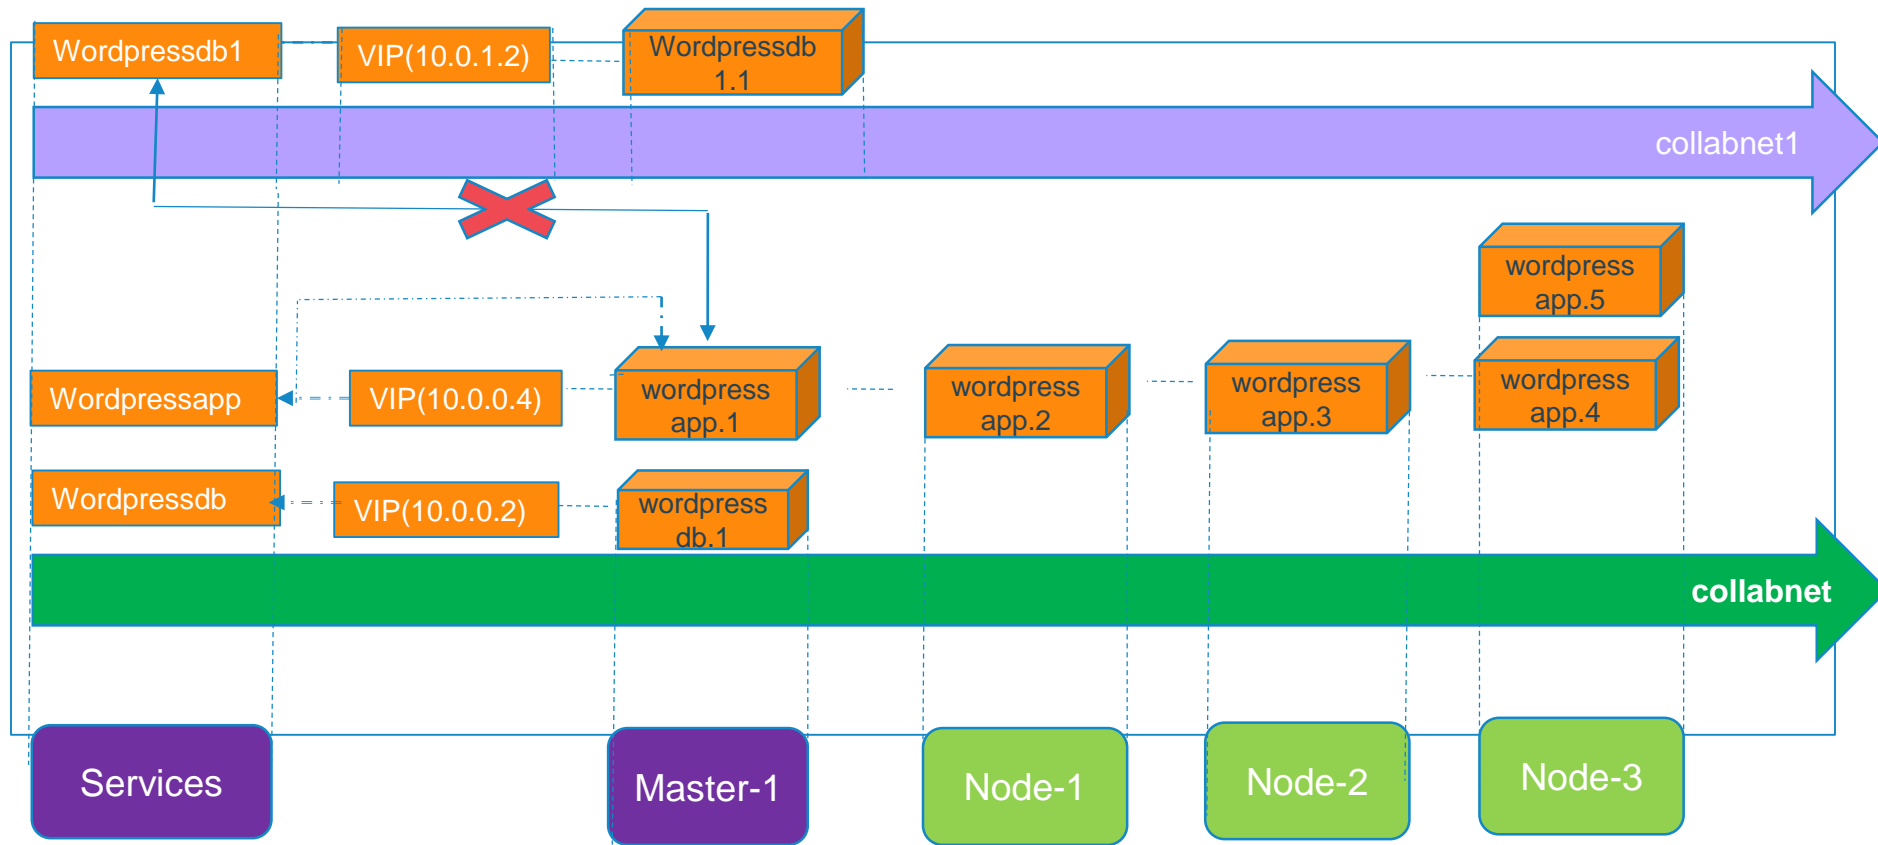
```
[{"NetworkID":"9eyjm4uv4ynmz0aubfqxise29","Addr":"10.0.0.2/24"}]
```

```
master==>sudo docker service inspect wordpressapp| tail -n30
        "Protocol": "tcp",
        "TargetPort": 80,
        "PublishedPort": 80
    }
    ],
    "Ports": [
        {
            "Protocol": "tcp",
            "TargetPort": 80,
            "PublishedPort": 80
        }
    ],
    "VirtualIPs": [
        {
            "NetworkID": "c4caizphmdpuhm1gjdle8eaal",
            "Addr": "10.255.0.7/16"
        },
        {
            "NetworkID": "9eyjm4uv4ynmz0aubfqxise29",
            "Addr": "10.0.0.4/24"
        }
    ]
    },
    "UpdateStatus": {
        "StartedAt": "0001-01-01T00:00:00Z",
        "CompletedAt": "0001-01-01T00:00:00Z"
    }
}
]
```

Verifying Service Discovery



Network – A Scope of Service Discoverability



Load Balancing



Load-Balancing

Distributes requests among the healthy nodes.

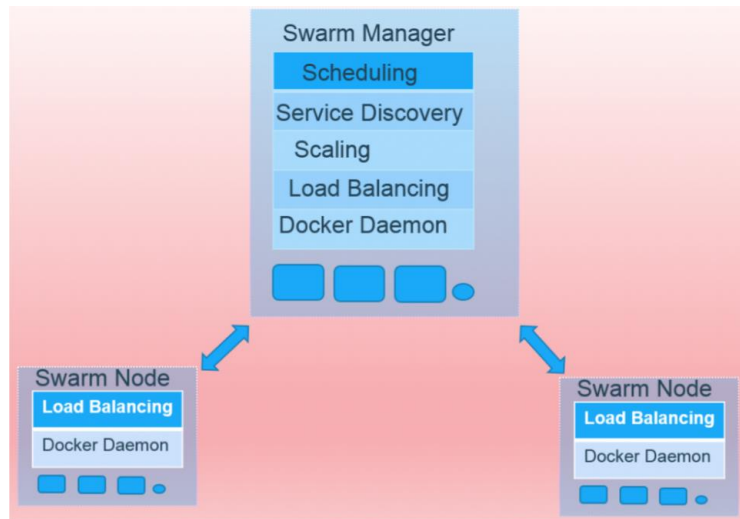
Decentralized, Highly Available – LB instance plumbed into every container instance

Internal Load Balancer – Provided by Embedded DNS

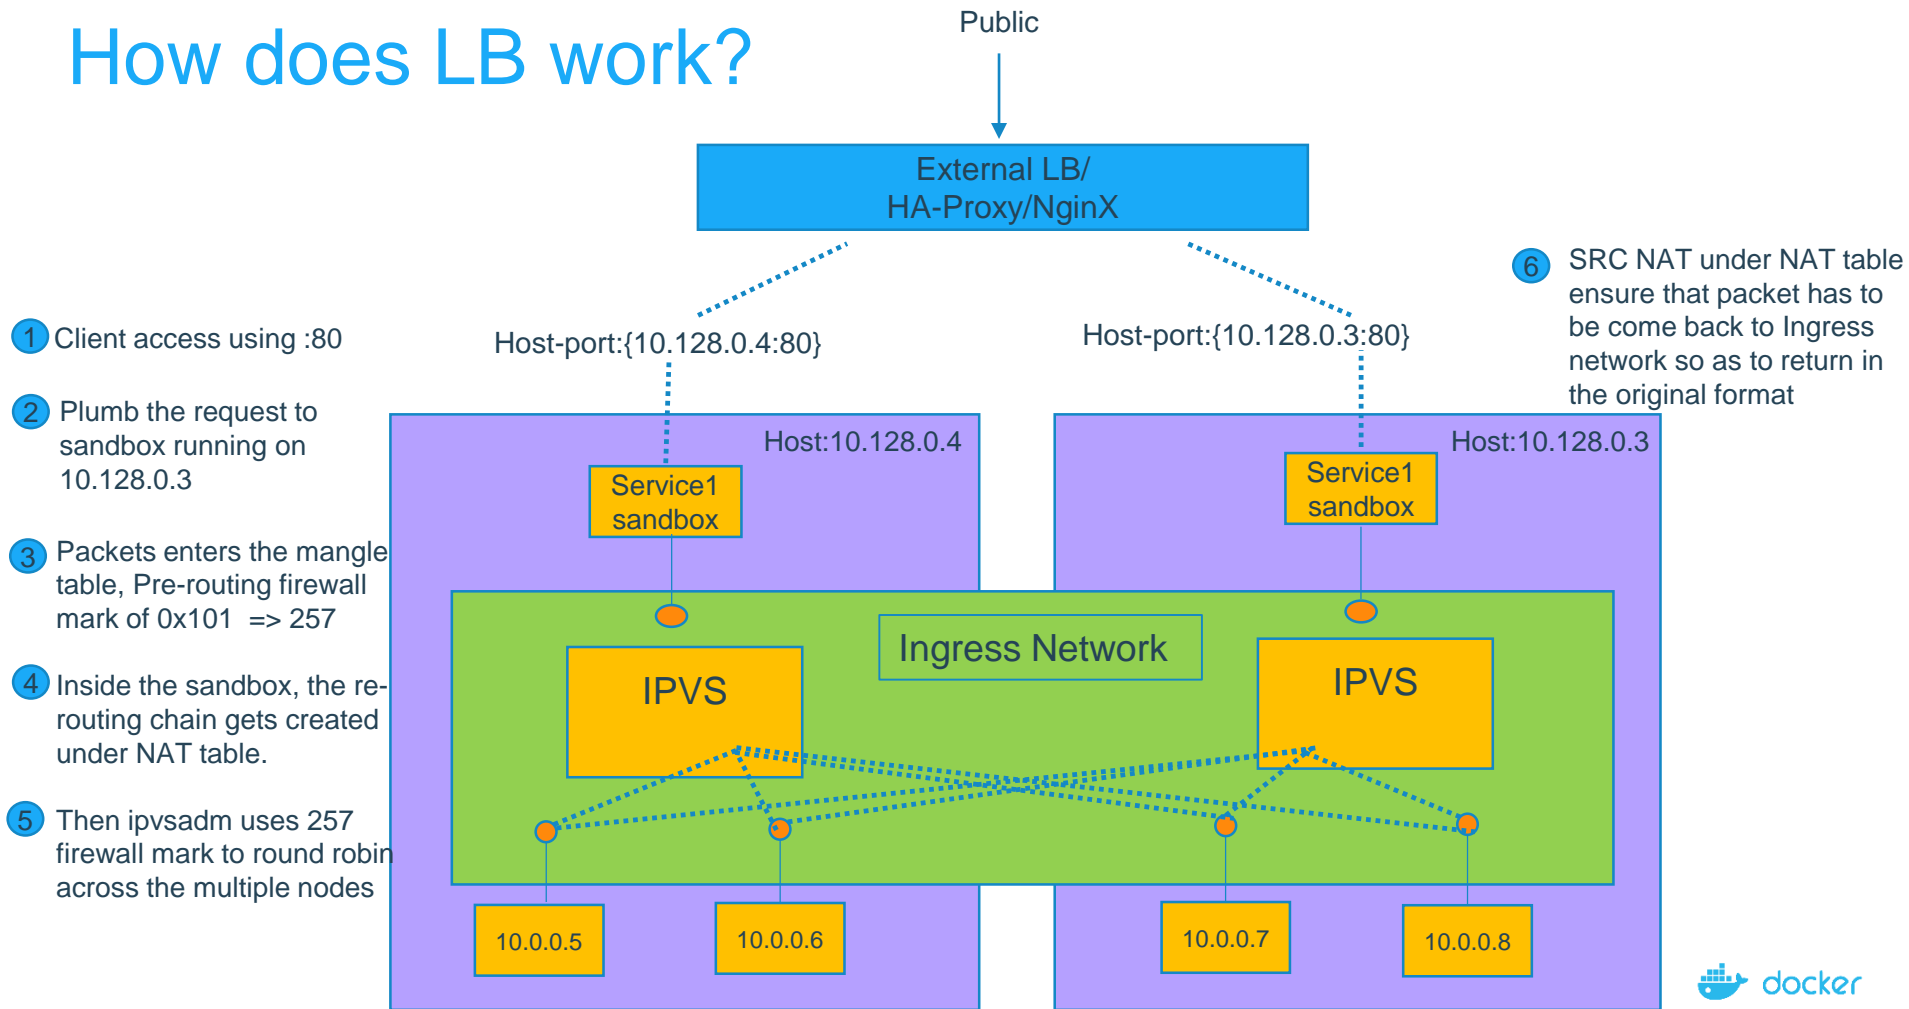
Can be used to discover both service & tasks

VIP based services uses IPVS(IP Virtual Server)

Kernel module (ip_vs) for LB



How does LB work?



Accessing the network sandbox

How to find the sandboxID?

```
root@master1:~# docker inspect f0cb | grep -i sandbox
    "SandboxID": "5fcd63cde83887d9033cac9fecdlf2c66de6b71648d660a6f80355be08aebb9f",
    "SandboxKey": "/var/run/docker/netns/5fcd63cde838",
root@master1:~# docker inspect d38c5 | grep -i sandbox
    "SandboxID": "36c249f1ef1dfa1fed03b85126588d4e3922dd5cb426d0c97ff800c5418e0dff",
    "SandboxKey": "/var/run/docker/netns/36c249f1ef1d",
root@master1:~# ls /var/run/docker/netns/
1-9eyjm4uv4y  1-c4caizphmd  36c249f1ef1d  5fcd63cde838  bf4eb2c9e6d7
```

Where's sandbox located?

```
root@master1:~# ls /var/run/docker/netns
1-9eyjm4uv4y  1-c4caizphmd  36c249f1ef1d  5fcd63cde838  bf4eb2c9e6d7
root@master1:~# nsenter --net=/var/run/docker/netns/5fcd63cde838 sh
```

Network namespace managed by overlay network driver(creating a bridge, terminating VXLAN tunnel etc.

Inspecting the sandbox

```
# iptables -t mangle -nvL
Chain PREROUTING (policy ACCEPT 291 packets, 45678 bytes)
  pkts bytes target    prot opt in     out     source    destination
    164 10845 MARK      tcp  --  *      *       0.0.0.0/0  0.0.0.0/0          tcp dpt:80 MARK set 0x101

Chain INPUT (policy ACCEPT 164 packets, 10958 bytes)
  pkts bytes target    prot opt in     out     source    destination

Chain FORWARD (policy ACCEPT 127 packets, 34720 bytes)
  pkts bytes target    prot opt in     out     source    destination

Chain OUTPUT (policy ACCEPT 164 packets, 10870 bytes)
  pkts bytes target    prot opt in     out     source    destination
      0      0 MARK      all  --  *      *       0.0.0.0/0  10.255.0.7          MARK set 0x101

Chain POSTROUTING (policy ACCEPT 291 packets, 45590 bytes)
  pkts bytes target    prot opt in     out     source    destination
```

```
root@master1:~# nsenter --net=/var/run/docker/netns/bf4eb2c9e6d7 sh
# iptables -t nat -nvL
Chain PREROUTING (policy ACCEPT 0 packets, 0 bytes)
  pkts bytes target    prot opt in     out     source    destination
    45  2160 REDIRECT tcp  --  *      *       0.0.0.0/0  0.0.0.0/0          tcp dpt:80 redir ports 80

Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
  pkts bytes target    prot opt in     out     source    destination

Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
  pkts bytes target    prot opt in     out     source    destination
      0      0 DOCKER_OUTPUT all  --  *      *       0.0.0.0/0  127.0.0.11

Chain POSTROUTING (policy ACCEPT 0 packets, 0 bytes)
  pkts bytes target    prot opt in     out     source    destination
      0      0 DOCKER_POSTROUTING all -- *      *       0.0.0.0/0  127.0.0.11
    45  2160 SNAT      all  --  *      *       0.0.0.0/0  0.0.0.0/0          ipvs to:10.255.0.3

Chain DOCKER_OUTPUT (1 references)
  pkts bytes target    prot opt in     out     source    destination
      0      0 DNAT      tcp  --  *      *       0.0.0.0/0  127.0.0.11          tcp dpt:53 to:127.0.0.11:33915
      0      0 DNAT      udp  --  *      *       0.0.0.0/0  127.0.0.11          udp dpt:53 to:127.0.0.11:56454

Chain DOCKER_POSTROUTING (1 references)
  pkts bytes target    prot opt in     out     source    destination
      0      0 SNAT      tcp  --  *      *       127.0.0.11  0.0.0.0/0          tcp spt:33915 to::53
      0      0 SNAT      udp  --  *      *       127.0.0.11  0.0.0.0/0          udp spt:56454 to::53
```

Routing Mesh

Routing Mesh is NOT Load-Balancer

Routing Mesh makes use of LB aspects

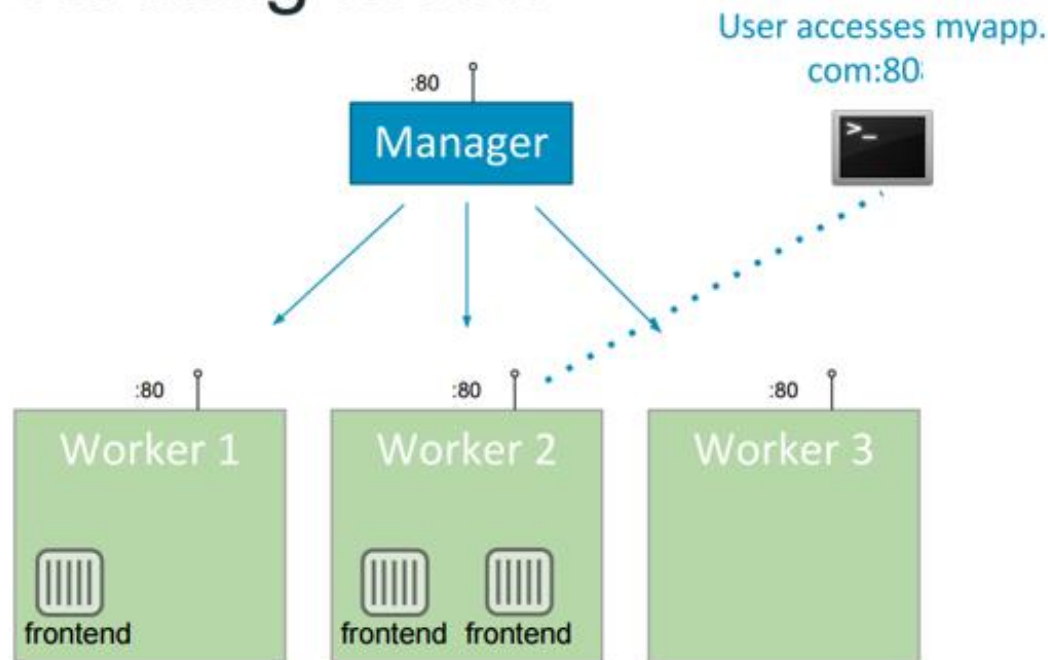
It provides global publish port for a given service

Built-in routing mesh for edge routing

Worker nodes themselves participate in ingress routing mesh

Port management at global Swarm Cluster level.

Routing Mesh



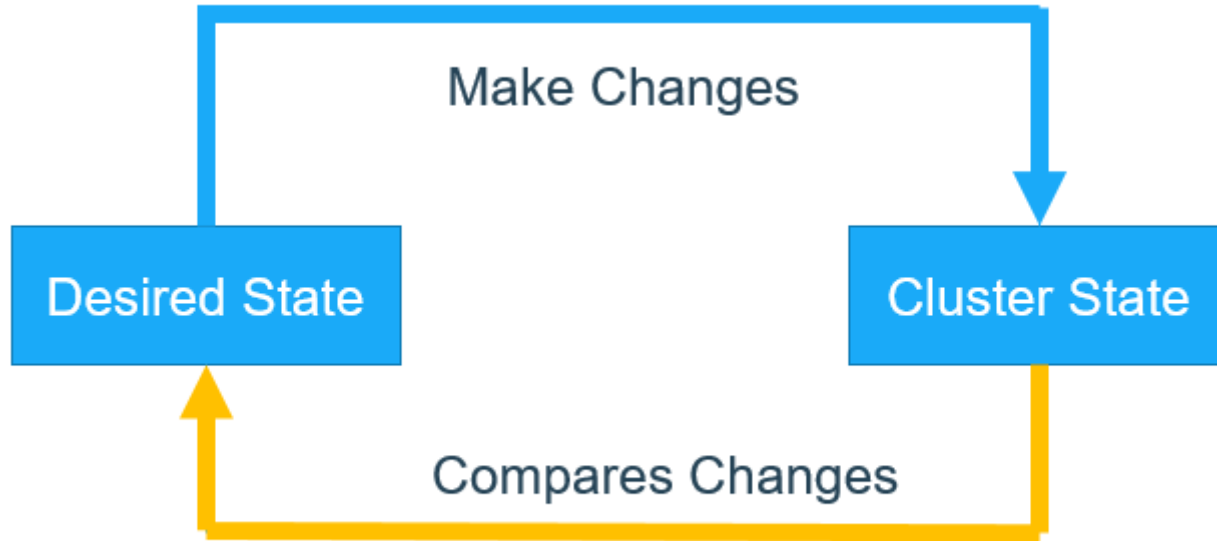
- Operator reserves a swarm-wide ingress port (80) for myapp
- Every node listens on 80
- Container-aware routing mesh can transparently reroute traffic from Worker3 to a node that is running container
- Built in load balancing into the Engine
- DNS-based service discovery



```
$ docker service create --replicas 3 --name frontend --network mynet  
--publish 80:80/tcp frontend_image:latest
```



Desired State Reconciliation




Building Our First Swarm Service



Hello world!


My hostname is 637fec5f215a



```
$docker network create -d overlay mynetwork
```


```
$docker service create --name mycloud --replicas 3 --network mynetwork --publish 80:80/tcp  
dockercloud/hello-world
```

Swarm Services – [Desired State \neq Actual State]




`$docker service scale mycloud=8`

Swarm Services – [Desired State = Actual State]




`$docker service scale mycloud=8`

Swarm Mode – Global Services




```
$docker service create --mode=global --name mycloud dockercloud/hello-world
```

Swarm Mode – High Availability Scheduling




Prioritizing spreading out the containers instead of equalizing the number of containers per node

Swarm Mode – High Availability Scheduling



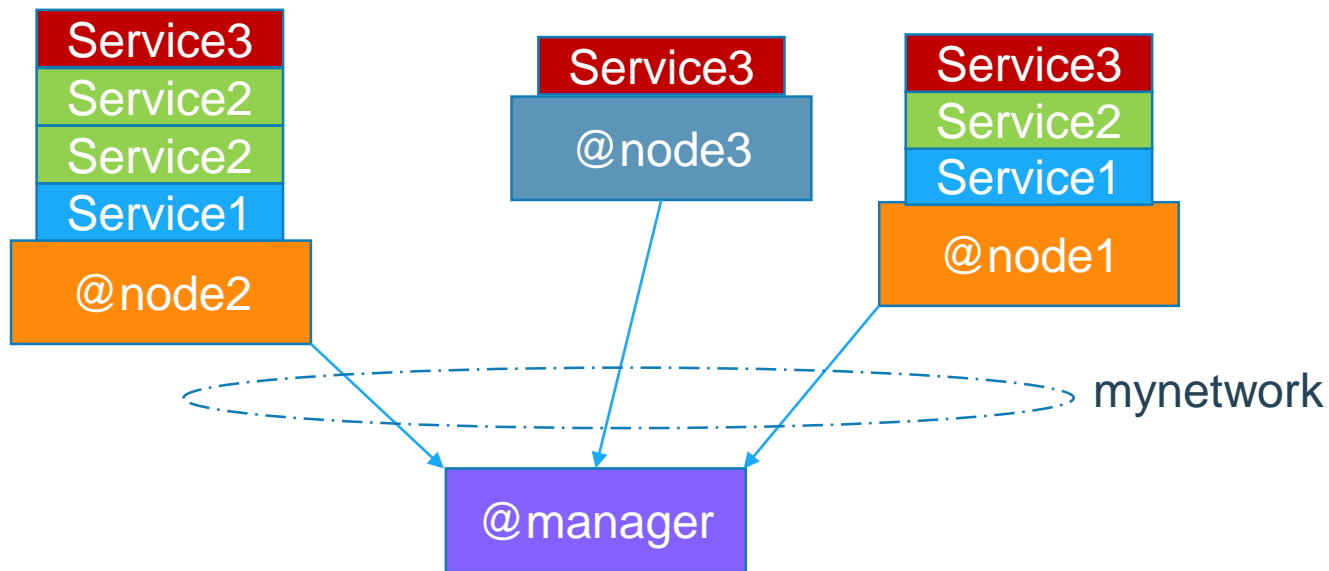
Adding a new Node - @node3

Swarm Mode – High Availability Scheduling




Prioritizing spreading out the containers instead of equalizing the number of containers per node

Swarm Mode – High Availability Scheduling




Prioritizing spreading out the containers instead of equalizing the number of containers per node

Swarm Mode – Topology Aware Scheduling




```
$docker node update --label-add datacenter=east node-1
```

Topology Aware Scheduling – How to use it?



```
$docker service create --replicas 2 --name wordpressdb1 \  
--network collabnet --placement-pref "spread=node.labels.datacenter" \  
--env MYSQL_ROOT_PASSWORD=collab123 \  
--env MYSQL_DATABASE=wordpress mysql:latest
```


Swarm Mode – Placement Constraints



```
$docker service create --network collabnet \  
  --endpoint-mode dnsrr \  
  --constraint 'node.platform.os == windows' \  
  --env ACCEPT_EULA=Y --env-file db-credentials.env \  
  --name db microsoft/mssql-server-windows
```

● node1
manager
31.405G RAM

● win00003S
manager
8.000G RAM

It's Demo Time

● myapp_web

image : python-web:latest@sha256:4b2adfe6a0125404cbd9dacee07a31e
tag : latest@sha256:d2215f98ba966
updated : 9/9 9:15
b2adfe6a0125404cbd9dacee07a31e
state : running

● myapp_redis

image : redis:3.2.100-nanoserver@sha256:81
tag : 3.2.100-nanoserver@sha256:81
updated : 9/9 9:18
20411388c259e10baba240ca24e93d
state : running

Welcome!

We're bypassing the Captcha and redirecting you now..

Start Session



docker con17^{EU}

Bella Center, Copenhagen

16-19 October





docker